

IETF Hackathon: Measurement & Analysis for Protocols Research Group (MAPRG)

IETF 105

20-21 July 2019

Montreal



Hackathon Plan

- **Problem:**

- **Related draft:** kIP: a Measured Approach to IPv6 Address Anonymization
<https://arxiv.org/abs/1707.03900>
<https://www.ietf.org/proceedings/99/slides/slides-99-maprg-kip-a-measured-approach-to-ipv6-address-anonymization-03.pdf>
<https://www.youtube.com/watch?v=qYtaKuzXaiM#t=59m55s>
- **Specific problem to solve:** Implement a modified PATRICIA trie or base-2 radix tree for “longest prefix match” making feasible for the tens to hundreds of billions of active IPv6 addresses used on the web today.

- **To solve it:** enhance aguri_tree to make portions of the tree *immutable*.
<https://www.ijlab.net/~kjc/software/#aguri>

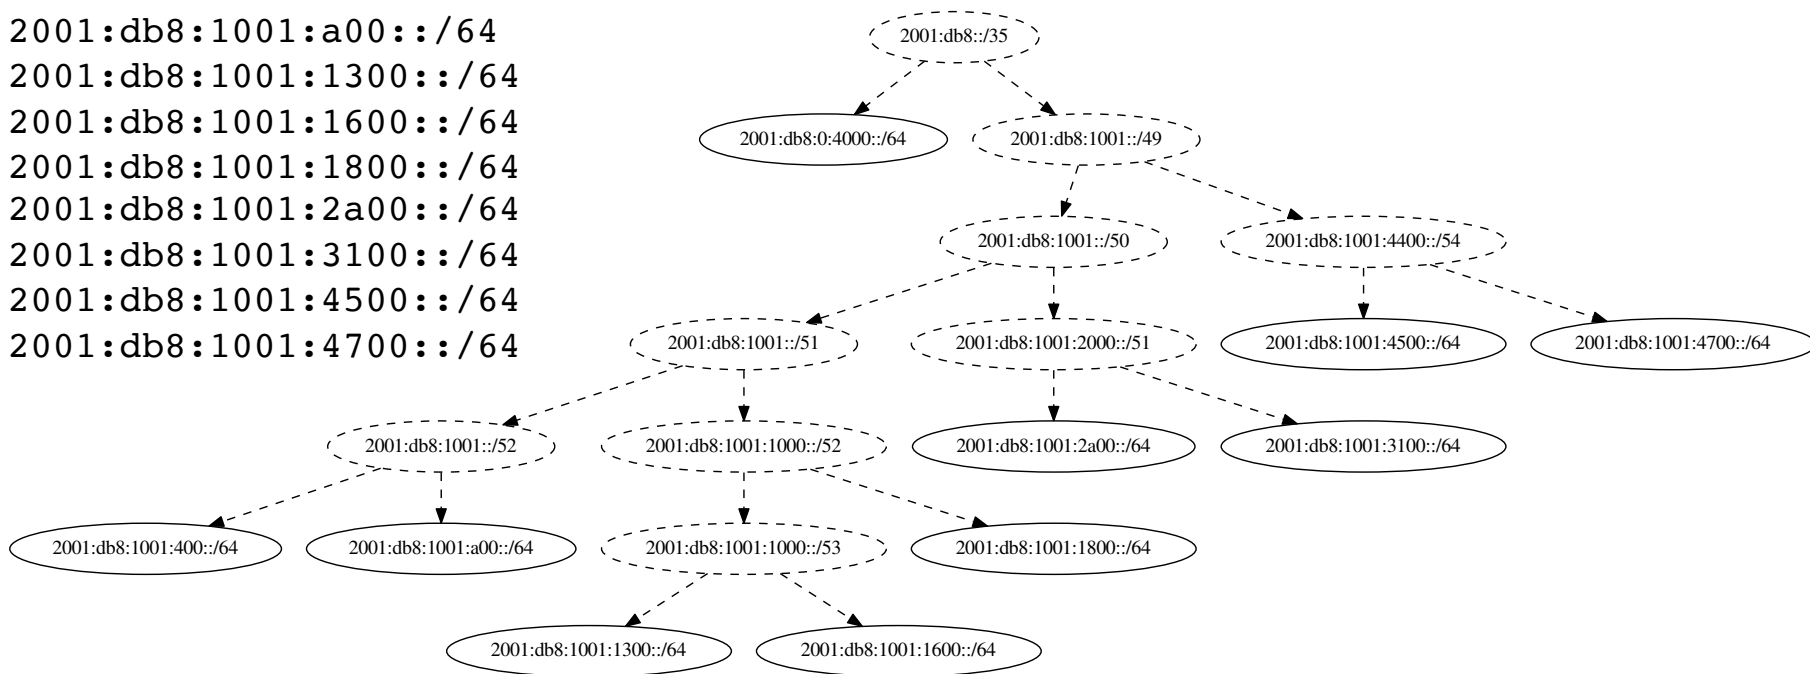
This allows partitioning of the problem for map-reduce/cluster operation, splitting the active IP addresses into manageable subsets (files) to produce intermediate results that are subsequently combined using prior mode (non-immutable nodes).

What got done

- <What you achieved? (key results)>
 - <New ideas - what team agreed on>
 - github pending
 - What was novel?
 - <Demos - links to videos>

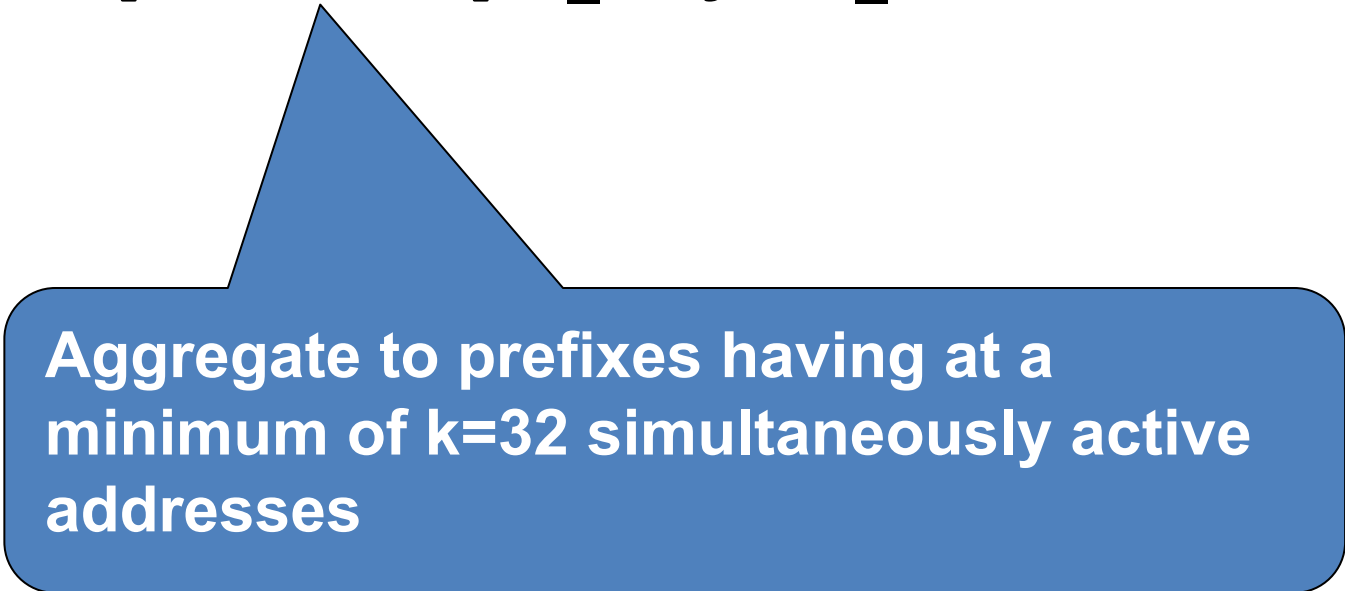
What got done

2001:db8:0:4000::/64
2001:db8:1001:400::/64
2001:db8:1001:a00::/64
2001:db8:1001:1300::/64
2001:db8:1001:1600::/64
2001:db8:1001:1800::/64
2001:db8:1001:2a00::/64
2001:db8:1001:3100::/64
2001:db8:1001:4500::/64
2001:db8:1001:4700::/64



What got done

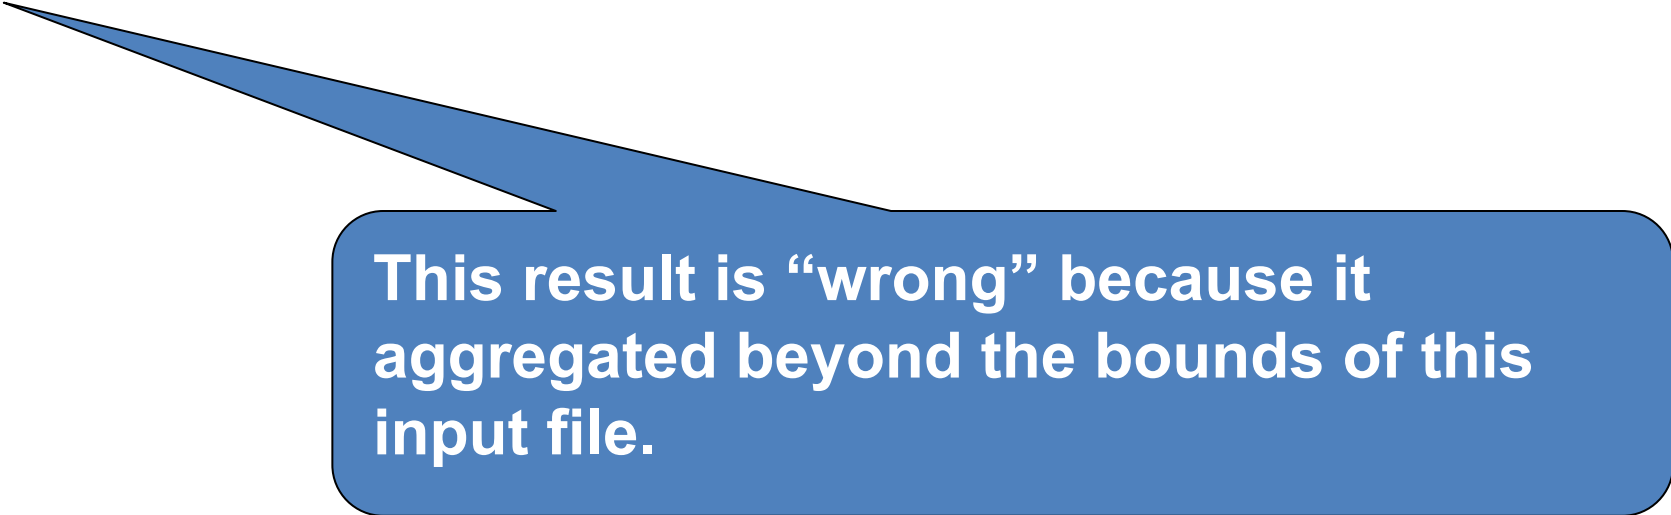
```
$ ./agurify -v -p50 -k32 input_fragment_v6.txt  
::/0
```



Aggregate to prefixes having at a minimum of k=32 simultaneously active addresses

What got done

```
$ ./agurify -v -p50 -k32 input_fragment_v6.txt  
::/0
```

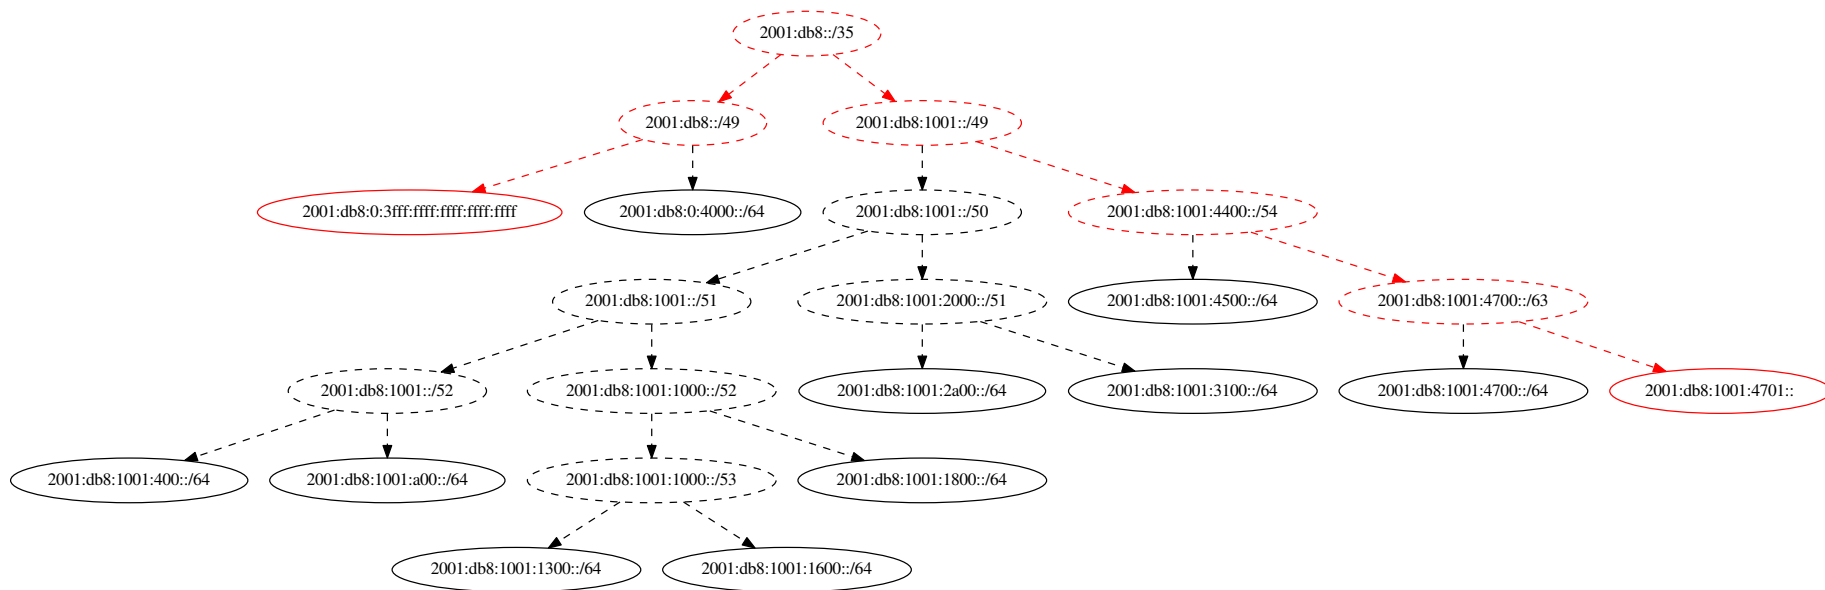


This result is “wrong” because it aggregated beyond the bounds of this input file.

What got done

2001:db8:0:3fff:ffff:ffff:ffff:ffff added immutable...

2001:db8:1001:4701:: added immutable...



What got done

```
$ ./agurify -v -i -p50 -k32 input_fragment_v6.txt  
20010db800003ffffffffff added immutable...  
20010db8100147ffffffffff added immutable...
```

Produce an intermediate result, i.e., don't aggregate to any prefix which could cover any address not appearing in this input file.

What got done

```
$ ./agurify -v -i -p50 -k32 input_fragment_v6.txt
20010db800003ffffffffffffffffffffffff added immutable...
20010db8100147010000000000000000 added immutable...
    2001:db8:0:4000::/64 335
    2001:db8:1001::/50 335
    2001:db8:1001:4500::/64 335
    2001:db8:1001:4700::/64 335
```

**This is a useful intermediate result
that can be processed, iteratively,
toward the final result.**

What we learned

- Lessons learned from this hackathon
 - Implemented IPv6 address “bignum” arithmetic rather than using GMP (GNU Multiple Precision Arithmetic Library)
 - New implementation/operation guidance?
 - Candidate best practice for aggregation-based IP address anonymization for privacy, *e.g.*, GDPR compliance
 - Aids investigation of address assignment practice, *e.g.*, to produce homogenous end user aggregates for matching with content in delivery networks
 - New work to take to WG?

A reference implementation of kIP

Wrap Up

Team members:
Dave Plonka

{based on code by
Kenjiro Cho,
Ryo Kaizaki}

MAPRG meets Friday morning:

<https://datatracker.ietf.org/rg/maprg/about/>
<https://trac.ietf.org/trac/irtf/wiki/map>

maprg-chairs@ietf.org